

# Parallel Algorithm Examples

Pangfeng Liu  
National Taiwan University

March 6, 2022

# Embarrassingly Parallel

- An embarrassingly parallel computation is a collection of tasks that require *none* or *little* communication among them. In other words, they are *independent*.
- This is “embarrassing” since nothing needs to be done to get good parallel performance.
- People doing parallel processing, e.g. me, are not fond of this kind of computation.

# Monte Carlo Method

- The Monte Carlo method repeats a *random process* to compute the answer.
- We generate random numbers as input to the computation to derive the answer from the random process. Note that all computations are *independent*.
- It is important to use different *random seed* with these tasks so that the results from them are *statistically independent*.

# Compute $\pi$

- Randomly throw darts into a square with an inscribed circle.
- Compute the number of darts that fall into and outside the circle.
- Compute the probability that a dart falls into the circle.
- Multiple the probability by 4 to approximate  $\pi$ .

# Notes

- This is not a good example because people will not compute  $\pi$  this way.
- Nevertheless this is an easy-to-understand example to illustrate the independence of tasks.

# Parameter Search

- Suppose we want to find a set of “best” parameters  $x = (x_1, \dots, x_n)$ , which maximize an objective function  $y = f(x)$ .
- We also assume that the  $f$  function is very complex so that we cannot deduce the values of  $f(x)$ 's for  $x$ 's that we have not yet computed the function values, from those function values that we have already computed.
- It is easy to dispatch the computation of  $f(x)$  to processors to speed up the parameter search, since there is no dependency between these computations.

# File Serving

- A web server serves static HTML files to clients.
- The clients' requests are independent, so the web server serves the files in parallel, maybe using multiple threads.
- If there are multiple web servers, they can also serve the files in parallel.

# Summary

- Embarrassingly parallel computation has good speedup and efficiency because the tasks are independent and do not require much communication.
- It is trivial to dispatch tasks to processors if they require roughly the same amount of time.
- It is non-trivial to dispatch tasks to processors if they require a different amount of time. In this case, we need *load balancing*.



# Discussion

- Give an example of embarrassingly parallel computation.

# Divide and Conquer

- Divide-and-conquer is a common parallel algorithm design technique.
- As in a sequential divide-and-conquer algorithm, the problem is first divided into sub-problems.
- Unlike a sequential divide-and-conquer algorithm, a parallel algorithm solves (conquer) the sub-problem *in parallel*.
- Some communication may be necessary since the sub-problems may depend on each other.
- Finally, we combine the answers from individual sub-problems into the final answer.

# Sequential Summation

- We want to sum  $n$  numbers, and  $n$  is huge.
- We add all numbers to the first one, and at the end, it has the sum.
- How many additions do this algorithm use?
- Is it possible to improve this algorithm?

## Lower and Upper Bounds

- The lower bound is the minimum amount of operations we need to solve this problem.
- For every algorithm  $A$ , there exists an input so that the cost is at least  $L(n)$  (as a function of the input size  $n$ ).
- The upper bound is the minimum number of operations we need to solve for all possible inputs.
- There exists an algorithm  $A^*$ , and for all inputs, the cost is at most  $U(n)$  (as a function of the input size  $n$ ).

# Discussion

- How do you find a lower bound?
- How do you find an upper bound?
- Is it possible that  $L(n)$  is asymptotically larger than  $U(n)$ ?

# Summation

- We want to sum  $n$  numbers, and  $n$  is very large.
- It is easy to see that we can apply divide-and-conquer technique to solve the problem in parallel with  $p$  processors.

# Parallel Summation Algorithm

- 1 Partition the numbers so that each processor has roughly  $n/p$  numbers.
- 2 Each processor computes the sum of assigned numbers.
- 3 A processor collects all the partial sums from other processors and computes the final sum.

# Analysis

- We assume that first step takes very little time. This is the case for shared memory model, but not necessarily true for distributed memory model.
- The second steps takes  $O(\frac{n}{p})$  times.
- The third steps takes  $O(p)$  times.

The time complexity is as follows.

$$O(\frac{n}{p} + p) \tag{1}$$



# Discussion

- How do we minimize the  $O(\frac{n}{p} + p)$  by choosing the right  $p$ ?

# Communication

- Having one processor collect all the answers is not efficient.
- We partition the processor into two groups. Every processor in the first group sends its answer to the corresponding processors in the second group.
- We repeatedly do this until we have only one processor left, who should have the final answer.

# Parallel Summation Algorithm

- 1 Partition the numbers so each processor has  $n/p$  numbers.
- 2 Each processor computes the sum of its numbers.
- 3 Use the recursive algorithm to compute the final sum.
- 4 This is similar to the *tree optimization* in synchronization.

# Analysis

- We assume that first step takes very little time.
- The second steps takes  $O(\frac{n}{p})$  time.
- The third steps takes  $O(\log p)$  time because the depth of a complete binary tree of  $n$  nodes is about  $O(\log n)$ .

The final time complexity is as follows.

$$O(\frac{n}{p} + \log p) \tag{2}$$

# Discussion

- Describe the difference between the previous two algorithms.

## Observation

- The first term ( $\frac{n}{p}$ ) is *computation*. We can *never* reduce this part.
- The second term ( $\log p$ ) is about *communication*. We try our best to reduce this part.
- If we increase  $p$ , the computation time decreases and the communication time increases. That means we have more workers to share the workload, but we need to communicate more among more workers.

## More Observations

- It is important to balance the load among processors. We want to send  $(\frac{n}{p})$  data to each processor for processing.
- A shared memory implementation is significantly easier than a distributed memory implementation.
- The recursive (or tree-like) communication pattern is much more complicated than a naive one, and it requires complicated synchronization.

# Analysis

Speedup

$$k = \frac{n}{\frac{n}{p} + \log p} \quad (3)$$

Efficiency

$$e = \frac{n}{n + p \log p} \quad (4)$$



## Choice of $p$

- What is the best  $p$  in terms of speedup?
- Set  $\frac{n}{p} = \log p$  and solve  $p = \frac{n}{\log n}$ .
- The minimum parallel execution time  $\Theta(\log n)$  is achieved when  $p = \Theta(\frac{n}{\log n})$ .
- If we set  $p = \sqrt{n}$ , then the time will be  $\Theta(\sqrt{n})$ , which is much larger than the optimal  $\Theta(\log n)$ .

# Discussion

- Use calculus to compute the optimal  $p$  value.
- Is this *theoretical* optimal  $p$  useful in practice?

## Discussion

- What will happen if we set  $p$  to  $n$ ?
- What is the communication time of this algorithm?
- What are the upper and lower bounds of the number of stages this algorithm requires?
- Can you improve the algorithm?

# Prefix Sum

Given  $n$  numbers  $(x_1, \dots, x_n)$ , we want to compute all prefix sums as follows.

$$s_k = \sum_{i=1}^k x_i \quad (5)$$

## Compact An Array

- The prefix sum has various applications.
- If there are zeros and non-zeros in an array  $A$  and we only wish to keep the non-zeros in a new array  $B$ , then we can do a prefix sum on another array  $P$  with 0 and 1 to determine the positions of non-zeros in  $B$ .

## Pack an Array

A	before packing	4	0	9	0	0	8	5	7
		1	0	1	0	0	1	1	1
P	0 and 1								

prefix sum

B	new index after packing	1	1	2	2	2	3	4	5
		4	9	8	5	7			

# Fibonacci's Numbers

- The prefix sum has various applications, and it is not limited to summation.
- We all know Fibonacci's numbers.

$$f_i = \begin{cases} 0 & i = 0 \\ 1 & i = 1 \\ f_{i-1} + f_{i-2} & i \geq 2 \end{cases} \quad (6)$$

# Fibonacci's Numbers

$$f_1 = f_1 \quad (7)$$

$$f_2 = f_0 + f_1 \quad (8)$$

$$\begin{pmatrix} f_1 \\ f_2 \end{pmatrix} = \begin{pmatrix} 0 & 1 \\ 1 & 1 \end{pmatrix} \begin{pmatrix} f_0 \\ f_1 \end{pmatrix} \quad (9)$$

$$\begin{pmatrix} f_2 \\ f_3 \end{pmatrix} = \begin{pmatrix} 0 & 1 \\ 1 & 1 \end{pmatrix} \begin{pmatrix} f_1 \\ f_2 \end{pmatrix} \quad (10)$$

$$= \begin{pmatrix} 0 & 1 \\ 1 & 1 \end{pmatrix} \begin{pmatrix} 0 & 1 \\ 1 & 1 \end{pmatrix} \begin{pmatrix} f_0 \\ f_1 \end{pmatrix} \quad (11)$$

$$\begin{pmatrix} f_i \\ f_{i+1} \end{pmatrix} = \begin{pmatrix} 0 & 1 \\ 1 & 1 \end{pmatrix}^i \begin{pmatrix} f_0 \\ f_1 \end{pmatrix} \quad (12)$$



# Prefix Product

Given  $n$  matrices  $(x_1, \dots, x_n)$ , we want to compute all prefix products as follows.

$$s_k = \prod_{i=1}^k x_i \quad (13)$$

# Discussion

- Give an example of prefix sum application.

# Parallel Prefix Sum Algorithm

- Use the  $k$ -th processor to compute  $s_k$ .

# Analysis

- A sequential algorithm can do this easily in  $O(n)$  time.
- We assume that we use one processor per data, so  $p = n$ .
- The  $k$ -th processor requires  $O(k)$  time.
- The parallel time is the maximum of all processor time, hence  $O(n)$ .
- The speedup is  $O(1)$ , and the efficiency is  $(\frac{1}{p})$ .
- Not very efficient.

# Discussion

- What is wrong with the previous algorithm?

# Parallel Prefix Sum Algorithm

- To avoid doing duplicated work, we again use the  $k$ -th processor to compute  $s_k$ , but we get the result from the  $k - 1$ -th processor.

$$s_k = s_{k-1} + x_k \quad (14)$$

# Analysis

- Now, we do not duplicate work, but we need to wait.
- The  $k$ -th processor cannot compute its sum before receiving  $s_{k-1}$ .
- The result will go ripple-like from the first to the last processor like a wave-front.
- The parallel time is the maximum of all processor time, hence  $O(n)$ .
- The speedup is  $O(1)$ , and the efficiency is  $(\frac{1}{n})$ .
- Again not very efficient.

# Discussion

- What is wrong with the previous algorithm?



# A Better Algorithm

- There are  $\log n$  stages.
- In the  $i$ -stage every element adds the element  $2^i$  to the left to itself.
  - In the first stage every element adds the element to its left to itself.
  - In the second stage every element adds the element two elements to its left to itself.

## Parallel Prefix Sum Algorithm

```
for  $i \leftarrow 0$  TO  $\log n - 1$  do  
  for  $k \leftarrow 2^i$  TO  $n$  do  
     $x[k] \leftarrow x[k - 2^i] + x[k]$  {The  $k$ -th processor receives a partial sum  
    from the  $k - 2^i$ -th processor.}  
  end for  
end for
```

## Parallel Prefix

2	1	7	4	2	3	1	5
2	3	8	11	6	5	4	6
2	3	10	14	14	16	10	11
2	3	10	14	16	19	20	25

# Analysis

- There are  $\log n$  steps, since  $p = n$  now.
- A processor does a sum and receives a message, so the time is  $O(1)$ .
- The total parallel time is  $O(\log n)$ , which is much better than  $O(n)$  in previous approaches.
- The speedup is  $O(\frac{n}{\log n})$ , and the efficiency is  $(\frac{1}{\log n})$ .

## Discussion

- Prove that the algorithm is correct.
- What is the possible problem with the previous algorithm?


# Improvement

- It is not practical to use one processor per data since the number of data is much more than the number of processors in practice.
- We will assume that  $n$  is much larger than  $p$ ; hence we need to partition the data among processors.
- Now, each processor will compute the prefix sum of its data first, then use the previous algorithm to “patch” things up.

## Improved Algorithm

- 1 Partition data among processors.
- 2 Each processor computes its prefix sum.
- 3 Use the previous algorithm to compute the prefix sum of the *last* elements from all processors.
- 4 Use the prefix sum from the last elements to patch up the answers.

## Parallel Prefix





# Analysis

- A sequential algorithm can do this easily in  $O(n)$  time.
- The first step does not take time.
- Both the second and the fourth step take  $O(\frac{n}{p})$  time.
- The third step takes  $O(\log p)$ , as discussed before.

$$T_p = O\left(\frac{n}{p} + \log p\right) \quad (15)$$

- Similar optimization can find good  $p$ .

# Final Notes

- What did we compute?

$$s_4 = (((x_1 + x_2) + x_3) + x_4) \quad (16)$$

$$s_4 = ((x_1 + x_2) + (x_3 + x_4)) \quad (17)$$

# Discussion

- What property the operation must have in order for this algorithm to work?
- Does “+” have this property?
- Does “maximum” have this property?
- Does matrix multiplication have this property?

# Sorting

- To sort keys between 1 and  $n$  in order.
- We try the bucket sort first.

# Bucket Sort

- We need  $b$  buckets.
- We need to know the range of keys, and we assume that the keys are *evenly* distributed in this range.
- We use an array element to record whether a key *appears* in the input.
- We do not compare!

# Bucket Sort

- 1 Scan the keys and record its appearance in the corresponding buckets.
- 2 Read the keys from the buckets.

# Bucket Sort

## Example 1: Bucket sort

```

1 void bucketsort(int array[], int n, int b)
2 {
3     int i, j = 0;
4     int *bucket = calloc(b + 1, sizeof(int));
5     for (i = 0; i < n; i++)
6         bucket[array[i]]++;
7     for (i = 0; i <= b; i++)
8         while(bucket[i]--)
9             array[j++] = i;
10 }
```

1

<sup>1</sup><http://www.eecs.ucf.edu/courses/cop3502h/spr2007/sorting3.pdf>



# Bucket Sort

array

2	1	7	4	2	3	1	5
---	---	---	---	---	---	---	---

bucket

2	2	1	1	1	0	1	0	0	0
1	2	3	4	5	6	7	8	9	10

array

1	1	2	2	3	4	5	7
---	---	---	---	---	---	---	---



# Analysis

- The first step takes  $O(n)$  time.
- The second step takes  $O(n + b)$  time.
- The total complexity is  $O(n + b)$ .

## Discussion

- Why the second step takes  $O(n + b)$ , not  $O(nb)$  time?

# Parallel Bucket Sort

- 1 Every processor has exactly one bucket.
- 2 Every processor scans all keys to record keys corresponding to its bucket.
- 3 Read the keys from the buckets.

# Bucket Sort

array

2	1	7	4	2	3	1	5
---	---	---	---	---	---	---	---

bucket

2	2	1	1	1	0	1	0	0	0
1	2	3	4	5	6	7	8	9	10

array

1	1	2	2	3	4	5	7
---	---	---	---	---	---	---	---

# Parallel Bucket Sort

- Every processor takes  $O(n)$  time just to scan data.
- Exactly how to “read the keys from the bucket”?

# Discussion

- What is wrong with the previous algorithm?

# Parallel Bucket Sort

- Every processor reads only  $\frac{n}{p}$  keys, and record the appearance into its own set of buckets.
- Now it only takes  $O(\frac{n}{p})$  time.


## Read Keys

How do we read the keys out in the second step?

- 1 Every processor remembers the number of keys it places in every bucket.
- 2 Each processor computes the number of keys that should be in its bucket.
- 3 Use the parallel prefix sum algorithm to know each bucket's starting position.
- 4 Read from the bucket.



# Bucket Sort



# Analysis

- 1 The first step takes  $O(\frac{n}{p})$  time.
- 2 The second step takes  $O(b)$  because each processor needs to add  $b$  numbers.
- 3 The third step takes  $O(\log p)$ .
- 4 the fourth step takes  $O(\frac{n}{p})$ .<sup>2</sup>

---

<sup>2</sup>Assuming the keys are evenly distributed among processors.

## Final Time Complexity

- 1 The “read to bucket” takes  $O(\frac{n}{p})$  time.
- 2 The “read from bucket” takes  $O(\frac{n}{p} + \log p + b)$  time.
- 3 The total time complexity is  $O(\frac{n}{p} + \log p + b)$ .
- 4 The speedup is  $O(\frac{\frac{n+b}{\frac{n}{p} + \log p + b}}{\frac{n}{\log n}})$ , which is  $O(\frac{n+b}{\log n + b})$  when we set  $p$  to  $\frac{n}{\log n}$ .

# Discussion

- What is wrong with the previous algorithm?

# Improvement

- The bucket sort uses (and wastes) a lot of memory.
- We will try a “quicker” sort that also uses the divide-and-conquer techniques.
- Again, we assume that the keys have an even distribution, and we know the range.

## Sequential “Quicker” Sort

Sort the keys recursively as follows.

- Partition keys into  $g$  groups according to  $g - 1$  pivots.
- Individually sort the keys in each group recursively.
- Concatenate all the keys from the  $g$  groups.

# Analysis

- The first step takes  $O(n)$  time.
- Let the time to sort  $n$  keys be  $T(n)$ .

$$T(n) = gT\left(\frac{n}{g}\right) + n \quad (18)$$

- It is easy to see that  $T(n) = O(n \log_g n) = O(n \log n)$ .

# Discussion

- Explain what will happen when there are only two groups in the “quicker” sort.



## Parallel Quicker Sort

Every processor manages a group, which will store a range of keys.

- 1 Every processor reads only  $\frac{n}{p}$  keys and puts them into the corresponding bucket.
- 2 Each processor sorts the keys in its bucket.
- 3 Read all keys from processors.

# Pact an Array

array	2	1	7	4	2	3	1	5
-------	---	---	---	---	---	---	---	---

group	2	1	1	3	2	1	1	2	2	3
-------	---	---	---	---	---	---	---	---	---	---

group	4	5				4	5			
-------	---	---	--	--	--	---	---	--	--	--

group	7				
-------	---	--	--	--	--

7				
---	--	--	--	--

group 

--	--	--	--	--

--	--	--	--	--

array	1	1	2	2	3	4	5	7
-------	---	---	---	---	---	---	---	---

# Analysis

- The first step takes  $O(\frac{n}{p})$  time, without considering the synchronization.
- The second step takes  $O(\frac{n}{p} \log \frac{n}{p})$ , assuming that the keys are evenly distributed.
- The third step takes  $O(\log p)$  from previous analysis on prefix sum.
- The total time is  $O(\frac{n}{p} \log \frac{n}{p} + \log p)$ .
- The speedup is  $O(\frac{n \log n}{\frac{n}{p} \log \frac{n}{p} + \log p})$ .

## Discussion

- What synchronization mechanism do we need for the first step in the previous algorithm?
- Why we did not need it in the previous parallel bucket sort?

# Exchange Sort

- We consider a recursive “exchange” sort that is more suitable for distributed memory multicomputers.
- This algorithm is very suitable for the hypercube.
- We do not require that the keys have a uniform distribution since we can argue that the performance is *statistically* acceptable.

# Exchange Sort

- Divide the processors into two groups of equal size.
- Each processor *exchanges* keys with its corresponding processor in the other group according to a pivot – smaller keys go to a group of processors, and bigger keys go to the other group.
- We recursively do the same for both groups.

# Exchange Sort

- 1 Find a pivot.
- 2 Divide the processors into two groups of equal size.
- 3 Each processor “exchanges” keys with its corresponding processor in the other group.
- 4 We recursively do the same for both groups.
- 5 Finally, each processor sorts its keys.

# Pivot

- How to find a pivot?
- This is like the argument sequential quicksort; we randomly pick one, which is good enough.
- Use a binomial trial to argue that the tree depth of a quick sort is bounded by  $O(\log p)$  with high probability.



# Analysis

- We focus on the number of “movements” as the cost, since basically no computation is involved.
- In each level of the exchange a processor exchanges at most  $\frac{n}{p}$  keys.
- From previous argument the depth of the tree is bounded by  $O(\log p)$ , so the cost of exchange is  $O(\frac{n}{p} \log p)$ .
- Finally each processor still needs to sort its key with  $O(\frac{n}{p} \log \frac{n}{p})$  time.
- The final complexity is  $O(\frac{n}{p}(\log p + \log \frac{n}{p})) = O(\frac{n}{p} \log n)$ .

# Discussion

- What is the theoretical speedup of this algorithm?
- Find out the definition of hypercube and why is this algorithm suitable for hypercube.

# Matrix Multiplication

- Multiply two  $n \times n$  matrices  $A$ , and  $B$  and place the result into  $C$ .

$$A \times B = C \quad (19)$$

- We assume that the matrix is dense.

# Sequential Matrix Multiplication

- For the interest of simplicity we use the standard  $O(n^3)$  algorithm, instead of the Strassen<sup>3</sup> algorithm.
- The time complexity is  $O(n^3)$ .

---

<sup>3</sup>[http://en.wikipedia.org/wiki/Strassen\\_algorithm](http://en.wikipedia.org/wiki/Strassen_algorithm)

# Parallel Matrix Multiplication

- We use  $p$  processors to compute the  $n^3$  elements in  $C$ .
- Each processor simply computes the answer and no communication among them is necessary for a shared memory implementation.

# Analysis

- Each processor computes  $\frac{n^2}{p}$  elements.
- Each elements takes  $O(n)$  time to computes.
- The parallel time is  $O(\frac{n^3}{p})$ .
- The speedup is  $\frac{n^3}{\frac{n^3}{p}} = p$ . It seems to an embarrassingly parallel computation and nothing can be improved.

## Discussion

- Why no communication among processors is necessary for a shared memory implementation of the previous algorithm during the computation stage?

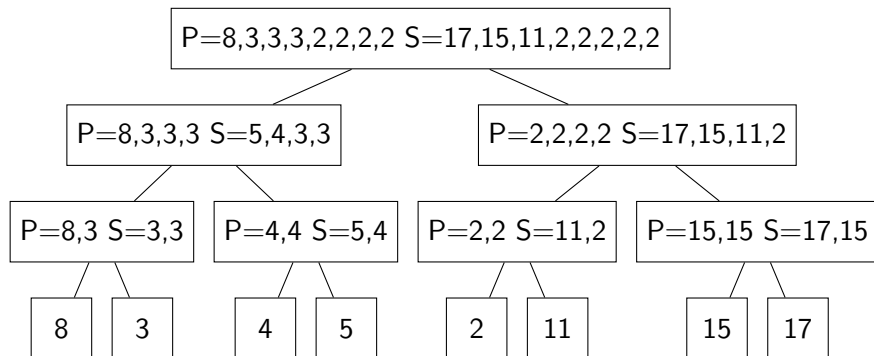
# Ranged Maxima

- Consider a sequence of  $n$  numbers.
- We want to preprocess these numbers with multiple processors, then answer a ranged maxima query in  $O(1)$  time.



# A Tree

- We build a tree from bottom to the top.
- Each leaf has a number.
  - Each internal node has two lists.
  - A  $P$  list has the prefix maxima for all the leaves in the subtree.
  - A  $S$  list has the suffix maxima for all the leaves in the subtree.



# Tree Construction

- How do we combine  $P$  and  $S$  from subtrees?
  - The first half of the  $P$  remains the same.
  - The first half of the  $S$  remains the same.
  - Compute the smaller value ( $x$ ) between the last elements of  $P$  and  $S$ .
  - Replace every value  $v$  in the second halves of  $P$  and  $S$  with  $\min(v, x)$ .

## Ranged Minima

- How do we find the ranged minimum  $(a, b)$ ?
- Find the least common ancestor  $r$  of  $a$  and  $b$ .
- Compute the minimum of the  $S$  element from the left subtree and the  $P$  element from the right subtree.

# Discussion

- Which elements from the left and the right subtree are we looking for?
- Why can we do it in  $O(1)$  time?

# Discussion

- How many processors do we need in the preprocessing?
- How much time do we need in the preprocessing?

# Discussion

- How to do this sequentially?
- What is the tradeoff between the preprocessing time and the query time?

# Graph

- A graph  $G$  consists of a node set  $V$  and an edge set  $E$ .
- For simplicity we assume that the edges are *undirected*.
- For simplicity the edge set  $E$  is an adjacency matrix.



# Degree

- The degree of a node is the sum of the *ones* in every row (or column).
- This is easily done in parallel using the parallel summation technique described earlier.
- One can verify whether there is a Euler cycle (path) in this graph by knowing the degrees.

# Breadth-First Search

- How to do BFS in parallel?
- Define  $T(v, k)$  to be the tree with  $v$  as the root and all nodes reachable from  $v$  by a path of length no more than  $2^k$ .

## Breadth-First Search

- $T(v, 0)$  is simply a node set  $v$  and its neighbors.
- For every node  $u$  in  $T(v, 0)$ , we merge  $T(u, 0)$  with  $T(v, 0)$  to get  $T(v, 1)$ .
- In general, for every node  $u$  in the frontier of  $T(v, i)$ , we merge  $T(u, i)$  with  $T(v, i)$  to get  $T(v, i + 1)$ .

## Discussion

- When can we stop? That is, what is the maximum value of  $k$  in  $v(v, k)$ ?
- How many processors do we need?